

# Singular Value Decomposition for Federated Data

***Claudia Serrano***<sup>1,2,3</sup> **and** ***Juan R. Gonzalez***<sup>\*1,2,3</sup>

<sup>1</sup>Autonomous University of Barcelona (UAB)

<sup>2</sup>ISGlobal - Barcelona Institute for Global Health

<sup>3</sup>Bioinformatics Research Group in Epidemiology (BRGE)

\*[juanr.gonzalez@isglobal.org](mailto:juanr.gonzalez@isglobal.org)

**2019-06-25**

Package

svdParallel 0.1.0

## Contents

1	Purpose of this package. . . . .	2
2	Getting started . . . . .	2
3	Computing a PCA using Jacobi algorithm . . . . .	3
4	Computing a PCA using Block method for DataSHIELD. . . . .	6
5	Computing a PCA using Block method for omic data. . . . .	10
	References . . . . .	17

## 1 Purpose of this package

---

Analyzing data across hospitals and institutions without the data leaving the hospitals and adding institutions to a trusted network is an important part of privacy preserving data analysis. In order to assure all of this we can use DataSHIELD, an open source solution that let researchers analyze microdata without physically sharing this data with them. It allows a standardized, monitored, indirect and secure access to data.

The aim of this package is to provide functions that compute the Singular Value Decomposition of large matrices and perform a Principal Component Analysis when we do not have access to the whole dataset, as in DataSHIELD <sup>1</sup>. At the same time, we are interested in the high computational cost that computing a SVD has, so we present some functions that calculate SVD's and PCA's in a more efficient way. The functions in this package can be classified in two different blocks: Jacobi algorithm and Block algorithm. The first group contains functions that perform a SVD using functions that apply the two sided Jacobi algorithm. This is an iterative algorithm that is also parallelizable and could perform well with large datasets and multi cores processors. However, it cannot be applied to DataSHIELD. The second group uses an incremental algorithm that we call Block method and consists of partitioning the matrix under study in submatrices in order to reduce their dimension and compute the SVD faster. This is applicable to DataSHIELD.

<sup>1</sup><http://www.datashield.ac.uk/>

The Block method allows us to compute a PCA of different datasets without being able to join them. However, to do so we need that each dataset contains the same variables and a greater number of individuals than variables. This usually happens in most of the datasets. Nevertheless this is not the case for omic data. In this case, Block functions will not be useful for DataSHIELD but they will be good to perform efficient calculations in large datasets.

## 2 Getting started

---

First, you can install `svdParallel` from Github by typing

```
devtools::install_github("isglobal-brge/svdParallel")
```

Then, you can load the required packages to reproduce this document.

```
library(svdParallel)
library(FactoMineR)
library(irlba)
library(doParallel)
library(devtools)
library(Biobase)
library(MultiDataSet)
library(made4)
library(brgedata)
library(parallel)
library(SummarizedExperiment)
library(microbenchmark)
library(BiocStyle)
```

In the following sections of this document we will present some examples in order to make clear how the different functions work. To do so we will use three datasets: `iris`, `breast` and `breastTCGA`. The `iris` dataset is an R dataset and the other two, `breast` and `breastTCGA`, are saved in the data folder of this project, so you can access and manipulate all of them.

### 3 Computing a PCA using Jacobi algorithm

In this case we will use the `iris` dataset. We can see the head of this dataset as follows

```
head(iris)
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1         5.1         3.5         1.4         0.2  setosa
2         4.9         3.0         1.4         0.2  setosa
3         4.7         3.2         1.3         0.2  setosa
4         4.6         3.1         1.5         0.2  setosa
5         5.0         3.6         1.4         0.2  setosa
6         5.4         3.9         1.7         0.4  setosa
```

We need to remove the last column because it is a categorical variable that we are not able to use in a PCA. We will use it later to analyze the results.

```
head(iris$Species)
[1] setosa setosa setosa setosa setosa setosa
Levels: setosa versicolor virginica

i <- iris[,-5]
```

The function that performs a PCA is `getPCA` and we have to set the method that we want to use to compute this PCA. The default method is the Block method, so if we want to compute the PCA using the Jacobi method we need to specify it. After computing a PCA with `getPCA` we can use function `print` to see the main values obtained with the PCA, the variance of each component and the percentage of variance explained.

```
gJ <- getPCA(i,method="Jacobi")
print(gJ)
**Results for the Principal Component Analysis using method 'Jacobi' **

Variance of the components
      PC1      PC2      PC3      PC4
[1,] 2.918498 0.9140305 0.1467569 0.02071484

Percentatge of variance explained for each component
      PC1      PC2      PC3      PC4
[1,] 72.96245 22.85076 3.668922 0.5178709
```

Now, we can plot the results using function `plot`. We need to specify what type of plot we want, individuals or variables. Moreover, if we do not specify anything more the plot will be made in the first two principal components.

```
par(mfrow=c(1,2))
plot(gJ, type="variables")
plot(gJ, type="individuals")
```

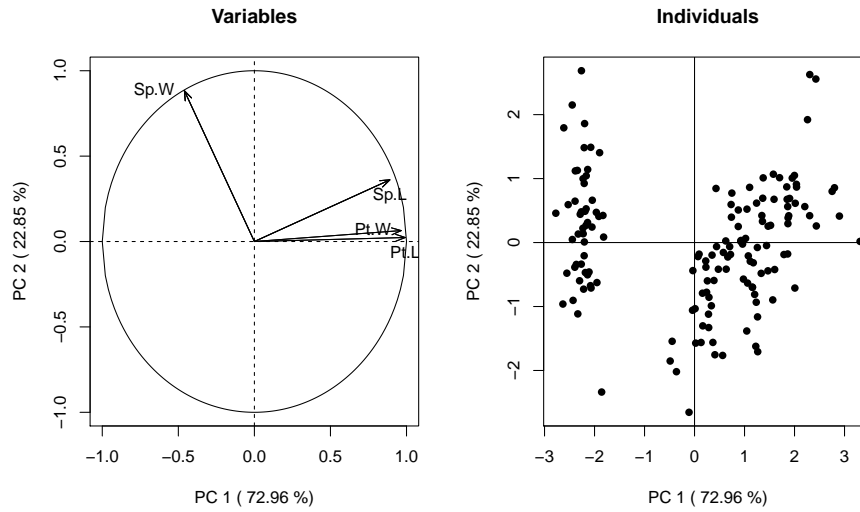


Figure 1: [Variables](#) and [individuals](#) projections from `getPCA` function using Jacobi method  
Left plot shows the variables in to the two first principal components, while right plot corresponds to the individuals.

If we want the plots in other components we can do it by adding `comps=c(i,j)`, where `i` and `j` are the new components.

```
par(mfrow=c(1,2))
plot(gJ, type="variables", comps=c(1,3))
plot(gJ, type="individuals", comps=c(1,3))
```




Figure 2: [Variables](#) and [individuals](#) projections from `getPCA` function using Jacobi method  
Left plot shows the variables into the first and third first principal components, while right plot corresponds to the individuals.

The `getPCA` function returns a list of different variables that we can use to perform a deeper analysis of the data. These variables are the coordinates of the variables, the Y matrix containing the coordinates of the individuals, the variance and percentage of variance explained by each component, the V matrix of the SVD performed and the method used.

In order to access to these variables we can do

```
gJ$varcoord
```

	PC1	PC2	PC3	PC4
Sepal.Length	0.8901688	0.36082989	0.27565767	-0.03760602
Sepal.Width	-0.4601427	0.88271627	-0.09361987	0.01777631
Petal.Length	0.9915552	0.02341519	-0.05444699	0.11534978
Petal.Width	0.9649790	0.06399985	-0.24298265	-0.07535950

For instance, we can use the Y matrix with the coordinates of the individuals to represent them according to the species they belong:

```
setosa <- iris$Species == "setosa"
versicolor <- iris$Species == "versicolor"
virginica <- iris$Species == "virginica"
plot(gJ$Y[setosa,c(1,2)], main="Individuals",
     type="p", col="blue", xlab="PC1 (72.96%)",
     ylab="PC2 (22.85%)", xlim=c(-4,5.5))
points(gJ$Y[versicolor,c(1,2)], col="red")
points(gJ$Y[virginica,c(1,2)], col="green")
legend("bottomright",
      c("Setosa", "Versicolor", "Virginica"),
      col=c("blue", "red", "green"), pch=1,
      cex=0.75)
```

As we can see, the individuals are classified according to the species.

We can perform the same calculations using `JacobiR` method. This method is not parallelizable and it is faster than `Jacobi`. We would recommend using `JacobiR` when we do not deal with big matrices and we do not need to parallelize. However, in small matrices the difference is not significant.

The functions that contain the Jacobi algorithm are `Jacobi`, `JacobiR` and `parJacobi` and they compute the eigenvalues and eigenvectors of a real symmetric matrix using the two-sided Jacobi algorithm. Functions `svdJacobi` and `svdJacobiR` compute a SVD of a real rectangular matrix and returns the left and right singular vectors and the singular values.

```
svdJ <- svdJacobi(i)
svdJ$d
```

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
95.959914	17.761034	3.460931	1.884826

We have obtained the singular values of the iris dataset.

As these functions use an iterative method we need to set a tolerance. The value set by default is the tolerance of the machine, but it can be changed setting the `tol` parameter.




Figure 3: [Principal components of iris dataset](#)  
 The plot shows the samples according to the species they belong to.

## 4 [Computing a PCA using Block method for DataSHIELD.](#)

In this case we will also use `getPCA` function but we will specify a different method. If we want to simulate that we are in a DataSHIELD experiment and we cannot join the different datasets we can use a list containing the datasets. To do so, we will use `breast` dataset. We need to remove the categorical variables to be able to work with the numerical ones.

```
breast[1:5,1:3]
  id diagnosis radius_mean
1  842302      M      17.99
2  842517      M      20.57
3 84300903      M      19.69
4 84348301      M      11.42
5 84358402      M      20.29
num_breast <- breast[,3:32]
dim(num_breast)
[1] 569 30
```

As we can see, the dimension of the dataset is  $569 \times 29$ . We can create three datasets taking the first 200 individuals, the individuals in rows between 201 and 400 and the last individuals from 401 to 569. These datasets have the same variables and have more individuals than variables. We can use the Block method to compute the SVD. This method computes a SVD of each of the subsets and then joins the results to compute a SVD again. In the following figure we can see a flowchart of the procedure:

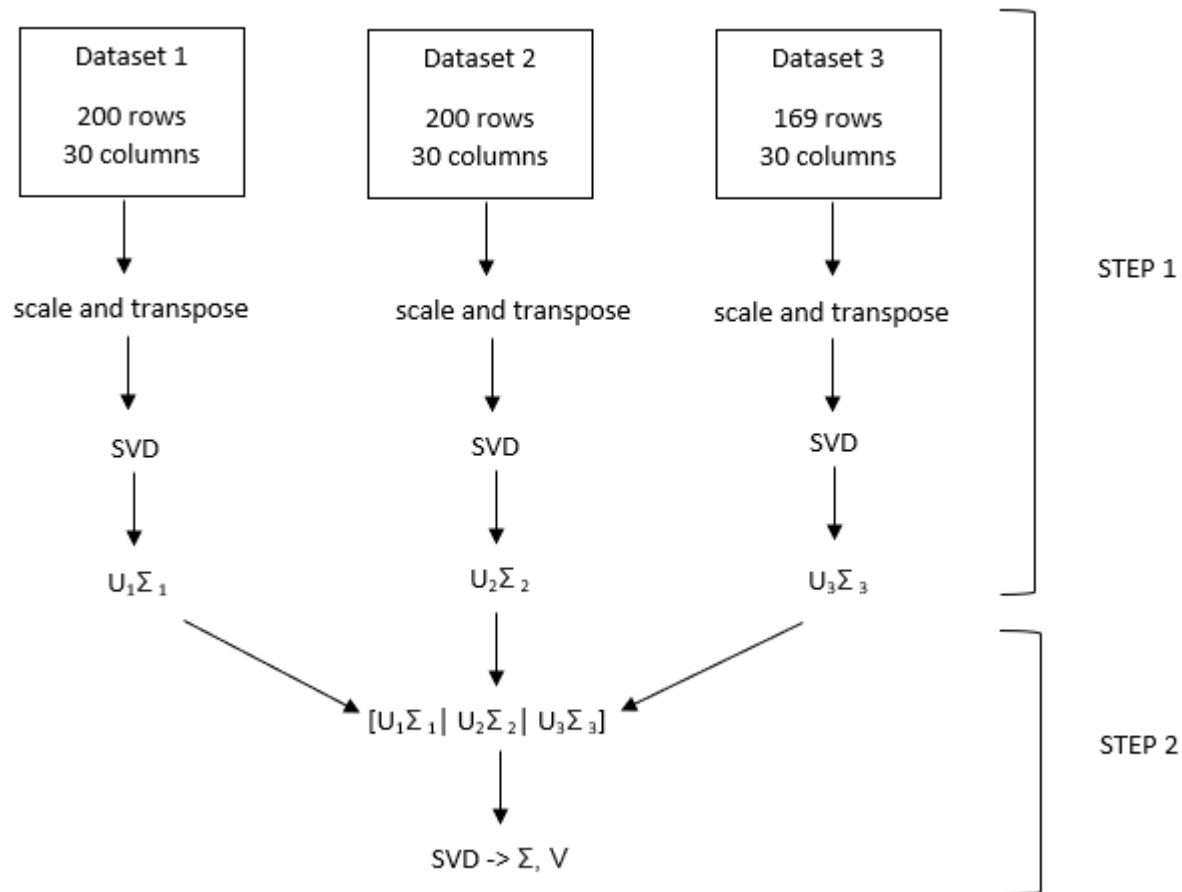


Figure 4: [Flowchart of the Block method procedure](#)

The SVD of each of the subdatasets is computed and after that, the results are joined to compute a SVD again.

To do so, we create a list containing these three datasets and we use `getPCA` function. In this case, we do not need to specify the method because the Block method is by default. It is also important that the `getPCA` function has the option to center and scale the data. By default the function will do it, but we can choose not to do so.

```
data1 <- num_breast[1:200,]
data2 <- num_breast[201:400,]
data3 <- num_breast[401:569,]

data_l <- list(data1,data2,data3)
```

```
gpcal <- getPCA(data_l)
```

Now we can plot the variables and the individuals again

```
par(mfrow=c(1,2))
plot(gpcal)
plot(gpcal, type="individuals")
```




Figure 5: Variables and individuals projections from `getPCA` function using Block method  
Left plot shows the variables into the two first principal components, while right plot corresponds to the individuals.

We can compare the results obtained with function `getPCA` with other functions in other R packages. IN the following plots we can see the results using `pca` function in `FactoMineR` package.

```
pcaR <- PCA(num_breast, graph=FALSE)
par(mfrow=c(1,2))
plot(pcaR, choix="var", label="none")
plot(pcaR, choix="ind", label="none")
```




Figure 6: Variables and individuals projections from `PCA` function  
Left plot shows the variables into the two first principal components, while right plot corresponds to the individuals.



The results do not match exactly. However, we can appreciate that these two plots are the two previous ones rotated 180°. This means that they are the same plots but using the opposite direction of the principal components. As we know, the principal components are not uniquely defined as we can take  $\mathbf{u}$  or  $-\mathbf{u}$  and they both would be a principal component. Therefore, the results are coherent.

Now we can also classify them according to the diagnosis of the tumor: benign or malignant.

```
m <- breast[,2] == "M"
b <- breast[,2] == "B"
plot(gpcaL$Y[m,c(1,2)], main="Individuals",
     type="p", col="blue", xlab="PC1 (44.27%)",
     ylab="PC2(18.97%)", xlim=c(-18,12),
     ylim=c(-14,9))
points(gpcaL$Y[b,c(1,2)], col="red")
abline(a=0,b=0, lty=3)
abline(v=0, lty=3)
legend("topright", c("Malignant", "Benign"),
      col=c("blue", "red"), pch=1, cex=0.75)
```




Figure 7: [Principal components of breast dataset](#)  
The plot shows the samples according to the group they belong (benign or malignant cancer)

We have been able to compute a PCA of three datasets together without merging them. This means that this function would be applicable to DataSHIELD experiments.

## 5 Computing a PCA using Block method for omic data.

In this case we will use a larger dataset with omic data. We will use a `SummarizedExperiment` called `breastMulti` that contains a `RangeSummarizedExperiment` called `breastTCGA`.

```
breastTCGA <- breastMulti[["expression"]]
breastTCGA
class: RangedSummarizedExperiment
dim: 8362 79
metadata(0):
assays(1): genexpr
rownames(8362): ABCA6 ABCC12 ... ZNF569 ZNF610
rowData names(0):
colnames(79): TCGA-C8-A12V TCGA-A2-A0ST ... TCGA-B6-A0X0
             TCGA-A2-A04Y
colData names(30): Gender Age.at.Initial.Pathologic.Diagnosis ...
                  Integrated.Clusters..unsup.exp. id
```

We can extract the dataset containing the genes information using `assay` function. There are 8362 variables and 79 individuals.

```
genes <- assay(breastTCGA)
dim(genes)
[1] 8362 79
genes[1:5,1:12]
```

	TCGA-C8-A12V	TCGA-A2-A0ST	TCGA-E2-A159	TCGA-BH-A0BW	TCGA-A2-A0SX
ABCA6	0.804100	1.411900	0.434300	-1.437600	0.907800
ABCC12	-0.653467	-0.411000	-0.499267	-0.939917	-0.622000
ABCC5	0.062833	-0.837000	-0.531000	-0.758667	0.034833
ABCG8	0.125750	-0.058250	-0.503250	-1.636250	0.003750
ABHD14B	-0.915313	-0.183188	-0.533312	-0.846313	-0.195687

	TCGA-AR-A1AI	TCGA-AR-A1AQ	TCGA-A1-A0S0	TCGA-A0-A0J6	TCGA-BH-A0B3
ABCA6	-0.280100	-0.084200	-0.893300	-1.996800	-0.563900
ABCC12	3.557667	-0.359833	-0.471167	-0.179267	-0.758500
ABCC5	0.731417	-1.078000	-0.435167	0.645667	-0.285167
ABCG8	0.343750	0.559750	0.240750	-0.891750	-0.852750
ABHD14B	-0.816063	-0.178312	-1.000063	-0.383188	-0.440437

	TCGA-C8-A134	TCGA-A2-A0T2
ABCA6	-0.075100	-1.545900
ABCC12	-0.321333	-0.616667
ABCC5	-0.591000	-0.384500
ABCG8	-0.456750	0.088750
ABHD14B	-0.717438	-0.059437

As this dataset contains omic data the individuals are in columns and the variables are in rows. This is not the normal way to build datasets and our functions require the individuals to be in the rows and the variables in the columns. This means that we need to apply `getPCA` to the transpose of the `genes` dataset. Moreover, the Block method has a parameter called `parts`. This parameter sets the number of subdatasets to be created from the dataset under

study. The default value is 2. If  $n$  is the number of rows and  $p$  is the number of columns of our dataset, this number `parts` cannot be greater than  $\max(n,p)$ . We can write to set `parts=5`.

```
gpca <- getPCA(t(genes), parts = 5, center=T, scale=T)
plot(gpca, type="individuals")
```




Figure 8: [Individuals projections from getPCA function using Block method](#)  
The plot shows the individuals into the two first principal components

In this case we will not plot the variables because there are too many and the resulting plot is not interpretable.

In the individuals plot we can see two differentiated groups. These groups correspond to the estrogen status, that can be positive or negative.

```
group<-as.factor(breastTCGA$ER.Status)
plot(gpca$Y[group=="Positive",c(1,2)],
     main="Individuals", type="p",
     col="blue",xlab="PC1 (17.84%)",
     ylab="PC2 (7.61%)", xlim=c(-60,70),
     ylim=c(-70,60))
points(gpca$Y[group=="Negative",c(1,2)], col="red")
```

```
abline(a=0,b=0, lty=3)
abline(v=0, lty=3)
legend("topright", c("Positive", "Negative"),
      col=c("blue", "red"), pch=1, cex=0.75)
```

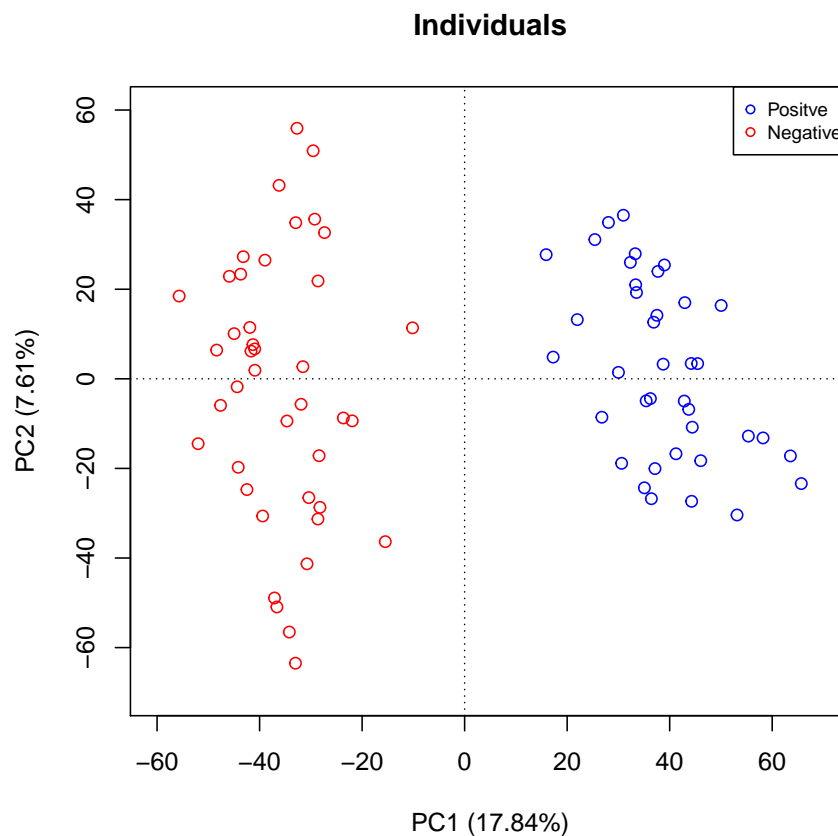


Figure 9: [Principal components of breastTCGA dataset](#)  
The plot shows the samples according to their estrogen status.

There exist specific functions to compute a PCA on gene expression data. This is the case of function `ord`. We can see that we obtain similar results.

```
group<-as.factor(breastTCGA$ER.Status)
out <- ord(genes, classvec=group, type = "pca")
plot(out, nlab=3)
plotarrays(out)
par(mfrow=c(1,1))
```

Apart from method `blockSVD` there is another method called `generalBlockSVD`. This method is the generalization from the previous one in a multi level algorithm. In order to use it we can set the parameters `k` and `q`. Their default values are 2 and 1, respectively. Parameter `q`

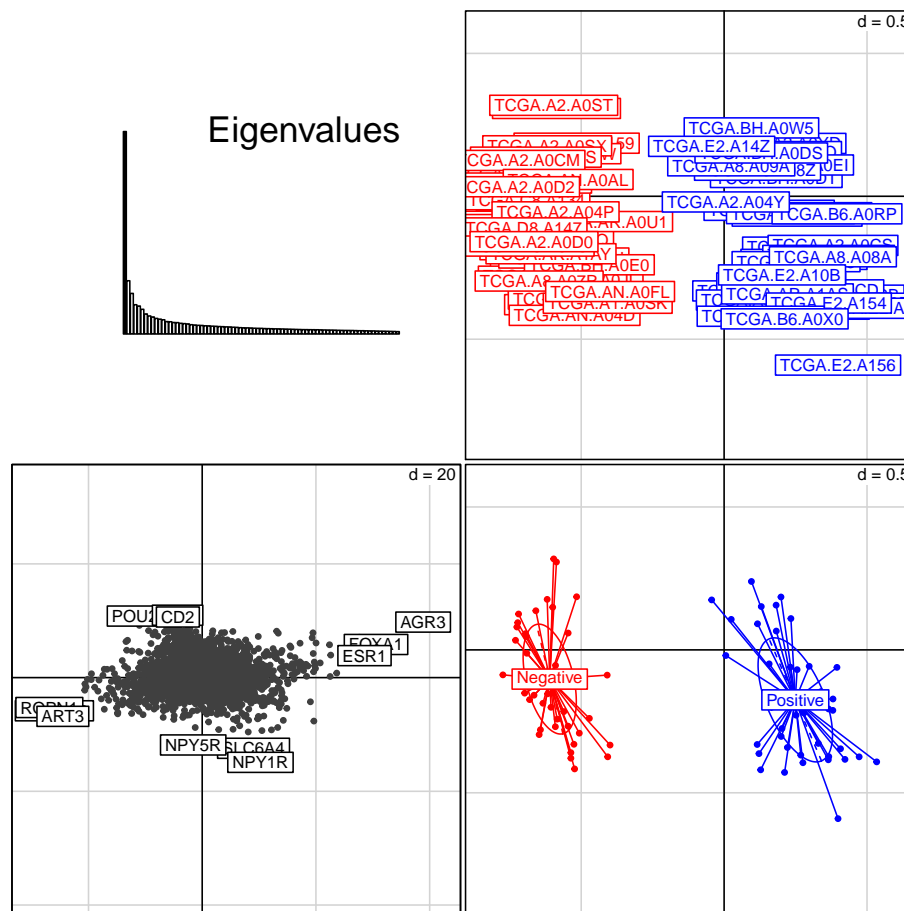


Figure 10: Variables and individuals projections from ord function  
First, there is the plot of the eigenvalues. Then, the plot of the variables (top right) and the individuals (bottom left) in the first two principal components, and the individuals classified in groups.

sets the number of levels of the algorithm (how many times we will repeat the process), and `k` sets the number of subdatasets to be concatenated at each level. More information about the method can be found at (Iwen and Ong 2016).

We can see that the same results can be obtained using this method:

```
gpca <- getPCA(t(genes), method="generalBlockSVD", k=2, q=2)
plot(gpca$Y[group=="Positive",c(1,2)],
     main="Individuals", type="p", col="blue",
     xlab="PC1 (17.84%)", ylab="PC2 (7.61%)",
     xlim=c(-60,70), ylim=c(-70,60))
points(gpca$Y[group=="Negative",c(1,2)], col="red")
abline(a=0,b=0, lty=3)
abline(v=0, lty=3)
legend("topright", c("Positive", "Negative"),
      col=c("blue", "red"), pch=1, cex=0.75)
```

Finally we can compare the different functions performance using the function `microbenchmark`. It computes the computation time of the functions and returns a table with the results.




Figure 11: [Individuals projections from `getPCA` function using General Block method](#)  
 The plot shows the individuals into the two first principal components classified according to their estrogen status.

```
mb <- microbenchmark("Block method (k=2)"= getPCA(t(genes),
  parts = 2, center=F, scale=F),
  "Block method (k=5)"=getPCA(t(genes),
  parts = 5, center=F, scale=F),
  "General Block method (k=2, q=2)"
= getPCA(t(genes), method="generalBlockSVD",
  k=2, q=2, center=F, scale=F),
  "General Block method (k=3, q=2)"
= getPCA(t(genes), method="generalBlockSVD",
  k=3, q=2, center=F, scale=F),
  "PCA function"=PCA(t(genes), graph=FALSE),
  "ord function" = ord(genes, classvec=group,
    type = "pca"), times=10)

autoplot.microbenchmark(mb)
```




Figure 12: Violinplot of the computational times for different functions

The two in the bottom are done using Block method and the following two are using General Block method. The second has been done using PCA function and the first one is using ord function.

If we look at the mean column we can see that the PCA function is the one that last more in average. This can be explained because this function computes a lot of different values a part from the necessary ones to compute a basic PCA. This also happens with function `ord`. These functions compute other scores and it makes it difficult to compare them with the others. The best way to analyze our functions' performance is to compare it to the function `getPCA0`. This function computes a basic PCA using only one SVD of the whole matrix. Our `getPCA` functions are faster in this case than PCA and `ord`, but it does not necessarily mean that they are more efficient.

```
mb2 <- microbenchmark("Block method (k=2)"= getPCA(t(genes),
  parts = 2, center=F, scale=F),
  "Block method (k=5)"= getPCA(t(genes),
  parts = 5, center=F, scale=F),
  "General Block method (k=2,q=2)"
  = getPCA(t(genes),
  method="generalBlockSVD", k=2, q=2,
  center=F, scale=F),
  "General Block method (k=3, q=2)"
```

```

= getPCA(t(genes),
method="generalBlockSVD", k=3, q=2,
center=F, scale=F),
"PCA0 no method" = getPCA0(t(genes),
center=F, scale=F), times=10)

autoplot.microbenchmark(mb2)

```




Figure 13: Violinplot of the computational times for different functions

The two in the bottom are done using Block method and the following two have been done using General Block method. The one in the top is using PCA0, it does not use any method.

The size of the `genes` dataset is not very big so the differences between the different functions are not significant. The performance of the Block method gets better when the size of the dataset increases.

When the size of the dataset is really large, we should use the parallel version of the two methods: `parBlockSVD` and `parGeneralBlockSVD`. They work as `blockSVD` and `generalBlockSVD`, respectively.



## References

Iwen, Mark, and Benjamin Ong. 2016. "A Distributed and Incremental Svd Algorithm for Agglomerative Data Analysis on Large Networks."